

Challenges – March 10, 2017 – Madrid, Spain

- Overview
 - Build a CM-managed CDH cluster and secure it
 - Place your work in the challenges/labs folder
 - All text files require Markdown (.md) extension and formatting
 - All screenshots must be in PNG format
 - You will create your own files for the challenges
 - You can consult with each other and research online
 - Submit only your own work!
 - Update your GitHub repo often -- don't wait until the end!
 - If you break your cluster, or your cluster breaks you:
 - Tell an instructor (mfernest or rsiwicki)
 - Review the work you have pushed to GitHub
 - Create a new Issue to describe what you think happened
-

Challenge Setup

- Create the Issue Challenges Setup
 - Make sure you have both mfernest and rsiwicki as Collaborators
 - Assign the Issue to yourself and label it started
 - In the file challenges/labs/0_setup.md:
 - List the cloud provider you are using (AWS, GCE, Azure, other)
 - List the nodes you are using by IP address and name
 - List the Linux release you are using
 - Demonstrate the disk capacity available on each node is ≥ 30 GB
 - List the command and output for `yum repolist` enabled
 - Add the following Linux accounts to all nodes
 - User neymar with a UID of 2010
 - User ronaldo with a UID of 2016
 - Create the group barca and add ronaldo to it
 - Create the group merengues and add neymar to it
 - List the /etc/passwd entries for neymar and ronaldo
 - Not the entire file!
 - List the /etc/group entries for barca and merengues
 - Not the entire file!
 - Push these updates to your GitHub repo
 - Label your Issue submitted
 - Assign the Issue to both instructors
-

Challenge 1: Install a MySQL server

- Create the Issue Install MySQL or Install MariaDB as appropriate
 - Assign the Issue to yourself and label it started
 - Install a MySQL 5.6 or MariaDB 5.5 server, as appropriate, on the first node listed in 0_setup.md
 - Use the appropriate YUM repository to install the package.
 - Copy the repo you're using to challenges/labs/1_my-database-server.repo.md
 - On all cluster nodes
 - Install the appropriate DB client package and JDBC connector jar
 - Start the database service
 - Create the following databases
 - scm
 - rman
 - hive
 - oozie
 - hue
 - sentry
 - Record the following in challenges/labs/1_db-server.md
 - The hostname of your db server node
 - The command and output for display your database server's version
 - The command and output for listing your created databases
 - Push this work to your GitHub repo
 - Label the Issue submitted and assign it to both instructors
-

Challenge 2: Install Cloudera Manager 5.9.x

- Create the Issue Install CM
 - Assign yourself to the Issue and label it started
 - Install Cloudera Manager on the second node listed in 0_setup.md
 - Configure the CM repo to install the latest release
 - List the command and output for `ls /etc/yum.repos.d` in `challenges/labs/2_cm.md`
 - Copy the `cloudera-manager.repo` file to `challenges/labs/2_cloudera-manager.repo.md`
 - Configure Cloudera Manager
 - Use the `scm_prepare_database.sh` script to write your `db.properties` file
 - List the full command line in `2_cm.md`
 - Start the Cloudera Manager server. Then in `challenges/labs/2_db.properties.md`:
 - Add the complete first line from your server log
 - Add the complete line that contains the phrase "Started Jetty server"
 - Add the full contents of your `db.properties` file
 - Push these changes to your GitHub repo and label the Issue 'submitted'
 - Assign the issue to both instructors
-

Challenge 3 – Install CDH 5.9.x

- Create the Issue Install CDH
 - Assign the issue to yourself and label it started
 - Install the latest CDH release; deploy Coreset services only
 - Name your cluster using your GitHub handle
 - Create user directories in HDFS for neymar and ronaldo
 - Add the following to 3_cm.md:
 - The command and output for `hdfs dfs -ls /user`
 - The output from the CM API call `../api/v14/hosts`
 - Login to Hue to install the Hive sample data
 - Capture a Hue screen that lists the Hive tables to `challenges/labs/3_hue_hive.png`
 - Push this work to your GitHub repo and label the Issue submitted
 - Assign the issue to both instructors
-

Challenge 4 – HDFS Testing

- Create the Issue Test HDFS
 - Assign the issue to yourself and label it started
 - As user neymar, use teragen to generate a 65,536,000-record dataset into eight files
 - Set the block size for this file to 16 MB
 - Name the target directory tgen640
 - Use the time command to capture job duration
 - Put the following in the file challenges/labs/4_teragen.md
 - The full teragen command and job output
 - The result of the time command
 - The command and output of hdfs dfs -ls /user/neymar/tgen640
 - The command and output to show how many blocks are stored under this directory
 - Push this work to your GitHub repo and label the Issue submitted
 - Assign the issue to both instructors
-

Challenge 5 – Kerberize the cluster

- Create the Issue Kerberize cluster
 - Assign the issue to yourself and label it started
 - Install an MIT KDC on the same node as the CM server
 - Name your realm after your GitHub handle
 - Use ES as a suffix
 - For example: RSIWICKI.ES
 - Create Kerberos principals for neymar, ronaldo, and cloudera-scm
 - Grant cloudera-scm the privileges needed to create principals and generate keytabs
 - Enable Kerberos for the cluster
 - Run the terasort program as neymar using the output target /user/neymar/tsort640m
 - Copy the command and full output to challenges/labs/5_terasort.md
 - Run the Hadoop pi program as the user ronaldo
 - Copy the command and full output to challenges/labs/5_pi.md
 - Copy the **text files** stored in /var/kerberos/krb5kdc/ to your repo:
 - Add the prefix 5_ and the suffix .md to the original file name
 - Example: 5_kdc.conf.md
 - Push this work to your GitHub repo and label the Issue submitted
 - Assign the issue to both instructors
-

Challenge 6 – Upgrade CM and CDH

- Create the Issue Update Cluster
 - Update Cloudera Manager to the latest available release
 - Capture the command used to determine the API version available, along with the output.
 - Store them in the file `6_cm_latest.png`
 - Upgrade CDH to the latest available release
 - Capture the CDH version and the services that are running
 - Store it in the file `6_cdh_latest.png`
 - Assign the issue to both instructors
 - Push all work to your GitHub repo
-

Once you finish, or when time is called:

- Commit any outstanding changes from your repo to GitHub
 - Notify mfernest@cloudera.com that you have stopped pushing to your repo
 - Continue working, if you like, after sending this note
 - Please fill out [this survey form \(https://goo.gl/forms/pmHeHx03zRu3cnlc2\)](https://goo.gl/forms/pmHeHx03zRu3cnlc2)
 - Anything else you'd like to express about the class?
 - Please add your comments to `labs/7_feedback_final.md` -- don't forget to commit them!
 - Now take it easy. You've worked very hard all week. Safe travels!
-